

Configure Spark Job Server on YARN

Deployment on YARN (Hortonworks 2.3.2)

Steps:

1. Download spark job server 0.6.0 that is compatible with Spark 1.4.1 installed in Hortonworks 2.3.3
2. Unzip and cd into spark-jobserver directory.
3. Copy config/local.sh.template to hdp.sh and edit as appropriate regarding DEPLOY_HOST, APP_USER, APP_GROUP, INSTALL_DIR and SPARK_HOME variables

Ex of hdp.sh:

```
# Environment and deploy file
# For use with bin/server_deploy, bin/server_package etc.
DEPLOY_HOSTS="sandbox.hortonworks.com"

APP_USER=root
APP_GROUP=hdfs
# optional SSH Key to login to deploy server
#SSH_KEY=/path/to/keyfile.pem
INSTALL_DIR=/usr/hdp/2.3.2.0-2950/spark-server
LOG_DIR=/var/log/job-server
PIDFILE=spark-jobserver.pid
JOBSERVER_MEMORY=1G
SPARK_VERSION=1.4.1
SPARK_HOME=/usr/hdp/2.3.2.0-2950/spark
SPARK_CONF_DIR=$SPARK_HOME/conf
# Only needed for Mesos deploys
SPARK_EXECUTOR_URI=/home/spark/spark-0.8.0.tar.gz
# Only needed for YARN running outside of the cluster
# You will need to COPY these files from your cluster to the remote machine
# Normally these are kept on the cluster in /etc/hadoop/conf
# YARN_CONF_DIR=/pathToRemoteConf/conf
# HADOOP_CONF_DIR=/pathToRemoteConf/conf
SCALA_VERSION=2.10.4
```

4. Copy config/local.conf.template to hdp.conf and edit as appropriate spark.master

Ex of hdp.conf

```
# Template for a Spark Job Server configuration file
# When deployed these settings are loaded when job server starts
#
# Spark Cluster / Job Server configuration
spark {
    # spark.master will be passed to each job's JobContext
    # master = "local[4]"
    # master = "mesos://vm28-hulk-pub:5050"
```

```

master = "yarn-client"

# Default # of CPUs for jobs to use for Spark standalone cluster
job-number-cpus = 4

jobserver {
  port = 8090
  jar-store-rootdir = /tmp/jobserver/jars

  jobdao = spark.jobserver.io.JobFileDAO

  filedao {
    rootdir = /tmp/spark-job-server/filedao/data
  }
}

# predefined Spark contexts
# contexts {
#   my-low-latency-context {
#     num-cpu-cores = 1           # Number of cores to allocate. Required.
#     memory-per-node = 512m      # Executor memory per node, -Xmx style eg 512m,
1G, etc.
#   }
#   # define additional contexts here
# }

# universal context configuration. These settings can be overridden, see README.md
context-settings {
  num-cpu-cores = 2           # Number of cores to allocate. Required.
  memory-per-node = 512m      # Executor memory per node, -Xmx style eg 512m, #1G,
etc.

  # in case spark distribution should be accessed from HDFS (as opposed to being
installed on every mesos slave)
  # spark.executor.uri = "hdfs://namenode:8020/apps/spark/spark.tgz"

  # uris of jars to be loaded into the classpath for this context. Uris is a string
list, or a string separated by commas ','
  # dependent-jar-uris =
["file:///some/path/present/in/each/mesos/slave/somepackage.jar"]

  # If you wish to pass any settings directly to the sparkConf as-is, add them here in
passthrough,
  # such as hadoop connection settings that don't use the "spark." prefix
  passthrough {
    #es.nodes = "192.1.1.1"
  }
}

# This needs to match SPARK_HOME for cluster SparkContexts to be created successfully
# home = "/home/spark/spark"
}

# Note that you can use this file to define settings not only for job server,
# but for your Spark jobs as well. Spark job configuration merges with this
configuration file as defaults.

```

```
# Timeout configuration
spray.can.server {
  idle-timeout = 210 s
  request-timeout = 200 s
}
```

5. Launch **bin/server_deploy.sh hdp** -- this packages the job server along with config files and pushes it to the remotes you have configured in hdp.sh

6. On the remote server, start it in the deployed directory with **server_start.sh** and stop it with **server_stop.sh** (test is port configured, 8090 by default, is in LISTENING mode)

Observation:

- for logging configuration one should copy **log4j-server.properties** from the installation directory to the remote machine

- for a better configuration one should read and configure file **application.conf** from job-server*.jar